

Section III - SELECT

3.1: Selecting all columns

```
SELECT * FROM CUSTOMER;
```

To limit the number of records returned, use a LIMIT. To limit the results to just 2 records:

```
SELECT * FROM CUSTOMER LIMIT 2;
```

3.2: Selecting specific columns

```
SELECT CUSTOMER_ID, NAME FROM CUSTOMER;
```

3.3: Expressions

First, select everything from PRODUCT

```
SELECT * FROM PRODUCT;
```

You can use expressions by declaring a TAXED_PRICE. This is not a column, but rather something that is calculated every time this query is executed.

```
SELECT PRODUCT_ID,  
DESCRIPTION,  
PRICE,  
PRICE * 1.07 AS TAXED_PRICE  
FROM PRODUCT;
```

In SQLiteStudio, you can hit CTRL + SPACE on Windows and Linux to show an autocomplete box with available fields. For Mac, you will need to enable that configuration in preferences.

You can also use aliases to declare an UNTAXED_PRICE column off the PRICE, without any expression.

```
SELECT PRODUCT_ID,  
DESCRIPTION,  
PRICE as UNTAXED_PRICE,  
PRICE * 1.07 AS TAXED_PRICE  
FROM PRODUCT;
```

SWITCH TO SLIDES FOR MATHEMATICAL OPERATORS

3.4: Using round() Function

```
SELECT PRODUCT_ID,  
DESCRIPTION,  
PRICE,  
round(PRICE * 1.07, 2) AS TAXED_PRICE  
  
FROM PRODUCT;
```

3.5: Text Concatenation

You can slap a dollar sign to our result using concatenation.

```
SELECT PRODUCT_ID,  
DESCRIPTION,  
PRICE AS UNTAXED_PRICE,  
'$' || round(PRICE * 1.07, 2) AS TAXED_PRICE  
FROM PRODUCT
```

You can merge text via concatenation. For instance, you can concatenate two fields and put a comma and space , in between.

```
SELECT NAME,  
CITY || ', ' || STATE AS LOCATION  
FROM CUSTOMER;
```

You can concatenate several fields to create an address.

```
SELECT NAME,  
STREET_ADDRESS || ' ' || CITY || ', ' || STATE || ' ' || ZIP AS SHIP_ADDRESS  
FROM CUSTOMER;
```

This works with any data types, like numbers, texts, and dates. Also note that some platforms use concat() function instead of double pipes ||

SWITCH TO SLIDES FOR EXERCISE

3.6: Comments

To make a comments in SQL, use commenting dashes or blocks:

```
-- this is a comment
```

```
/*  
This is a  
multiline comment  
*/
```

Section IV- WHERE

4.1: Getting year 2010 records

```
SELECT * FROM station_data
WHERE year = 2010;
```

4.2: Getting non-2010 records

```
SELECT * FROM station_data
WHERE year != 2010;

SELECT * FROM station_data
WHERE year <> 2010;
```

4.3: Getting records between 2005 and 2010

```
SELECT * FROM station_data
WHERE year BETWEEN 2005 AND 2010
```

4.4: Using AND

```
SELECT * FROM station_data
WHERE year >= 2005 AND year <= 2010
```

4.5: Exclusive Range

This will get the years between 2005 and 2010, but exclude 2005 and 2010

```
SELECT * FROM station_data
WHERE year > 2005 AND year < 2010
```

4.6: Using OR

```
SELECT * FROM station_data
WHERE MONTH = 3
OR MONTH = 6
OR MONTH = 9
OR MONTH = 12
```

4.7: Using IN

```
SELECT * FROM station_data
WHERE MONTH IN (3,6,9,12);
```

4.8: Using NOT IN

```
SELECT * FROM station_data
WHERE MONTH NOT IN (3,6,9,12);
```

4.9: Using Modulus

The modulus will perform division but return the remainder. So a remainder of 0 means the two numbers divide evenly.

```
SELECT * FROM station_data
WHERE MONTH % 3 = 0;
```

4.10: Using WHERE on TEXT

```
SELECT * FROM station_data
WHERE report_code = '513A63'
```

4.11: Using IN with text

```
SELECT * FROM station_data
WHERE report_code IN ('513A63','1F8A7B','EF616A')
```

4.12: Using length() function

```
SELECT * FROM station_data
WHERE length(report_code) != 6
```

4.13A: Using LIKE for any characters

```
SELECT * FROM station_data
WHERE report_code LIKE 'A%';
```

4.13B: Using Regular Expressions

If you are familiar with regular expressions, you can use those to identify and qualify text patterns.

```
SELECT * FROM STATION_DATA
WHERE report_code REGEXP '^A.*$'
```

4.14: Using LIKE for one character

```
SELECT * FROM station_data
WHERE report_code LIKE 'B_C%';
```

For LIKE, % is used in a different context than modulus %

4.15: True Booleans 1

```
SELECT * FROM station_data
WHERE tornado = 1 AND hail = 1;
```

4.16: True Booleans 2

```
SELECT * FROM station_data
WHERE tornado AND hail
```

4.17: False Booleans 1

```
SELECT * FROM station_data
WHERE tornado = 0 AND hail = 1;
```

4.18: False Booleans 2

```
SELECT * FROM station_data
WHERE NOT tornado AND hail;
```

4.19: Handling NULL

A NULL is an absent value. It is not zero, empty text ‘’, or any value. It is blank.

To check for a null value:

```
SELECT * FROM station_data
WHERE snow_depth IS NULL;
```

4.20: Handling NULL in conditions

Nulls will not qualify with any condition that doesn't explicitly handle it.

```
SELECT * FROM station_data
WHERE precipitation <= 0.5;
```

If you want to include nulls, do this:

```
SELECT * FROM station_data
WHERE precipitation IS NULL OR precipitation <= 0.5;
```

You can also use a `coalesce()` function to turn a null value into a default value, if it indeed is null.

This will treat all null values as a 0.

```
SELECT * FROM station_data
WHERE coalesce(precipitation, 0) <= 0.5;
```

4.21: Combining AND and OR

Querying for sleet or snow

Problematic. What belongs to the AND and what belongs to the OR?

```
SELECT * FROM station_data
WHERE rain = 1 AND temperature <= 32
OR snow_depth > 0;
```

You must group up the sleet condition in parenthesis so it is treated as one unit.

```
SELECT * FROM station_data
WHERE (rain = 1 AND temperature <= 32)
OR snow_depth > 0;
```

Section V- GROUP BY and ORDER BY

5.1: Getting a count of records

```
SELECT count(*) as record_count FROM station_data
```

5.2 Getting a count of records with a condition

```
SELECT count(*) as record_count FROM station_data
WHERE tornado = 1
```

5.3 Getting a count by year

```
SELECT year, count(*) as record_count
FROM station_data
WHERE tornado = 1
GROUP BY year
```

5.4 Getting a count by year, month

```
SELECT year, month, count(*) as record_count
FROM station_data
WHERE tornado = 1
GROUP BY year, month
```

5.5 Getting a count by year, month with ordinal index

```
SELECT year, month, count(*) as record_count
FROM station_data
WHERE tornado = 1
GROUP BY 1, 2
```

5.6 Using ORDER BY

```
SELECT year, month, count(*) as record_count
FROM station_data
WHERE tornado = 1
GROUP BY year, month
```

5.7 Using ORDER BY with DESC

```
SELECT year, month, count(*) as record_count
FROM station_data
WHERE tornado = 1
GROUP BY year, month
ORDER BY year DESC, month
```

5.8 Counting non-null values

```
SELECT COUNT(snow_depth) as recorded_snow_depth_count
FROM station_data
```

5.9 Average temperature by month since year 2000

```
SELECT month, AVG(temperature) as avg_temp
FROM station_data
WHERE year >= 2000
GROUP BY month
```

5.10 Average temperature (with rounding) by month since year 2000

```
SELECT month, round(AVG(temperature),2) as avg_temp
FROM station_data
WHERE year >= 2000
GROUP BY month
```

5.11 Sum of snow depth

```
SELECT year, SUM(snow_depth) as total_snow
FROM station_data
WHERE year >= 2005
GROUP BY year
```

5.12 Multiple aggregations

```
SELECT year,
SUM(snow_depth) as total_snow,
SUM(precipitation) as total_precipitation,
MAX(precipitation) as max_precipitation

FROM station_data
WHERE year >= 2005
GROUP BY year
```

EXERCISES

Flip to slides

5.13 Using HAVING

You cannot use WHERE on aggregations. This will result in an error.

```
SELECT year,
SUM(precipitation) as total_precipitation
FROM station_data
```



```
WHERE total_precipitation > 30
GROUP BY year
```

You can however, use HAVING.

```
SELECT year,
SUM(precipitation) as total_precipitation
FROM station_data
GROUP BY year
HAVING total_precipitation > 30
```

Note that some platforms like Oracle do not support aliasing in GROUP BY and HAVING.

Therefore you have to rewrite the entire expression each time

```
SELECT year,
SUM(precipitation) as total_precipitation
FROM station_data
GROUP BY year
HAVING SUM(precipitation) > 30
```

5.14 Getting Distinct values

You can get DISTINCT values for one or more columns

```
SELECT DISTINCT station_number FROM station_data
```

You can also get distinct combinations of values for multiple columns

```
SELECT DISTINCT station_number, year FROM station_data
```

Section VI - CASE Statements

6.1 Categorizing Wind Speed

You can use a CASE statement to turn a column value into another value based on conditions. For instance, we can turn different wind_speed ranges into HIGH, MODERATE, and LOW categories.

```
SELECT report_code, year, month, day, wind_speed,
CASE
  WHEN wind_speed >= 40 THEN 'HIGH'
  WHEN wind_speed >= 30 AND wind_speed < 40 THEN 'MODERATE'
  ELSE 'LOW' END
AS wind_severity
```

```
FROM station_data
```

6.2 More Efficient Way To Categorize Wind Speed

We can actually omit `AND wind_speed < 40` from the previous example because each `WHEN/THEN` is evaluated from top-to-bottom. The first one it finds to be true is the one it will go with, and stop evaluating subsequent conditions.

```
SELECT report_code, year, month, day, wind_speed,

CASE
    WHEN wind_speed >= 40 THEN 'HIGH'
    WHEN wind_speed >= 30 THEN 'MODERATE'
    ELSE 'LOW' END
as wind_severity

FROM station_data
```

6.3 Using CASE with GROUP BY

We can use `GROUP BY` in conjunction with a `CASE` statement to slice data in more ways, such as getting the record count by `wind_severity` and `year`.

```
SELECT year,

CASE
    WHEN wind_speed >= 40 THEN 'HIGH'
    WHEN wind_speed >= 30 THEN 'MODERATE'
    ELSE 'LOW' END
as wind_severity,

COUNT(*) as record_count

FROM station_data
GROUP BY 1,2
```

6.4 “Zero/Null” Case Trick

There is really no way to create multiple aggregations with different conditions unless you know a trick with the `CASE` statement. If you want to find two total precipitation, with and without tornado precipitations, for each year and month, you have to do separate queries.

Tornado Precipitation

```

SELECT year, month,
SUM(precipitation) as tornado_precipitation
FROM station_data
WHERE tornado = 1
GROUP BY year, month

```

Non-Tornado Precipitation

```

SELECT year, month,
SUM(precipitation) as non_tornado_precipitation
FROM station_data
WHERE tornado = 0
GROUP BY year, month

```

But you can use a single query using a CASE statement that sets a value to 0 if the condition is not met. That way it will not impact the sum.

```

SELECT year, month,
SUM(CASE WHEN tornado = 1 THEN precipitation ELSE 0 END) as tornado_precipitation,
SUM(CASE WHEN tornado = 0 THEN precipitation ELSE 0 END) as non_tornado_precipitation
FROM station_data GROUP BY year, month

```

Many folks who are not aware of the zero/null case trick will resort to derived tables (not covered in this class but covered in *Advanced SQL for Data Analysis*), which adds an unnecessary amount of effort and mess.

```

SELECT t.year,
t.month,
t.tornado_precipitation,
non_t.non_tornado_precipitation

FROM (
    SELECT year, month,
    SUM(precipitation) as tornado_precipitation
    FROM station_data
    WHERE tornado = 1
    GROUP BY year, month
) t

INNER JOIN

(
    SELECT year, month,
    SUM(precipitation) as non_tornado_precipitation
    FROM station_data
    WHERE tornado = 0
    GROUP BY year, month
) non_t

```

6.5 Using Null in a CASE to conditionalize MIN/MAX

Since NULL is ignored in SUM, MIN, MAX, and other aggregate functions, you can use it in a CASE statement to conditionally control whether or not a value should be included in that aggregation.

For instance, we can split up max precipitation when a tornado was present vs not present.

```
SELECT year,  
MAX(CASE WHEN tornado = 0 THEN precipitation ELSE NULL END) as max_non_tornado_precipitation,  
MAX(CASE WHEN tornado = 1 THEN precipitation ELSE NULL END) as max_tornado_precipitation  
FROM station_data  
GROUP BY year
```

Switch to slides for exercise

Exercise 6.1

SELECT the report_code, year, quarter, and temperature, where a “quarter” is “Q1”, “Q2”, “Q3”, or “Q4” reflecting months 1-3, 4-6, 7-9, and 10-12 respectively.

ANSWER:

```
SELECT  
  
report_code,  
year,  
  
CASE  
    WHEN month BETWEEN 1 and 3 THEN 'Q1'  
    WHEN month BETWEEN 4 and 6 THEN 'Q2'  
    WHEN month BETWEEN 7 and 9 THEN 'Q3'  
    WHEN month BETWEEN 10 and 12 THEN 'Q4'  
END as quarter,  
  
temperature  
  
FROM STATION_DATA
```

Exercise 6.2

Get the average temperature by quarter and month, where a “quarter” is “Q1”, “Q2”, “Q3”, or “Q4” reflecting months 1-3, 4-6, 7-9, and 10-12 respectively.

ANSWER

```

SELECT
year,

CASE
    WHEN month BETWEEN 1 and 3 THEN 'Q1'
    WHEN month BETWEEN 4 and 6 THEN 'Q2'
    WHEN month BETWEEN 7 and 9 THEN 'Q3'
    WHEN month BETWEEN 10 and 12 THEN 'Q4'
END as quarter,

AVG(temperature) as avg_temp

FROM STATION_DATA
GROUP BY 1,2

```

Section VII - JOIN

7.1A INNER JOIN

(Refer to slides Section VII)

View customer address information with each order by joining tables CUSTOMER and CUSTOMER_ORDER.

```

SELECT ORDER_ID,
CUSTOMER.CUSTOMER_ID,
ORDER_DATE,
SHIP_DATE,
NAME,
STREET_ADDRESS,
CITY,
STATE,
ZIP,
PRODUCT_ID,
ORDER_QTY

FROM CUSTOMER INNER JOIN CUSTOMER_ORDER
ON CUSTOMER.CUSTOMER_ID = CUSTOMER_ORDER.CUSTOMER_ID

```

Joins allow us to keep stored data normalized and simple, but we can get more descriptive views of our data by using joins.

Notice how two customers are omitted since they don't have any orders (refer to slides).

7.2B A BAD APPROACH

You may come across a style of joining where commas are used to select the needed tables, and a `WHERE` defines the join condition as shown below:

```
SELECT ORDER_ID,  
CUSTOMER.CUSTOMER_ID,  
ORDER_DATE,  
SHIP_DATE,  
NAME,  
STREET_ADDRESS,  
CITY,  
STATE,  
ZIP,  
PRODUCT_ID,  
ORDER_QTY  
  
FROM CUSTOMER, CUSTOMER_ORDER  
WHERE CUSTOMER.CUSTOMER_ID = CUSTOMER_ORDER.CUSTOMER_ID
```

Do not use this approach no matter how much your colleagues use it (and educate them not to use it either). It is extremely inefficient as it will generate a cartesian product across both tables (every possible combination of records between both), and then filter it based on the `WHERE`. It does not work with `LEFT JOIN` either, which we will look at shortly.

Using the `INNER JOIN` with an `ON` condition avoids the cartesian product and is more efficient. Therefore, always use that approach.

7.2 LEFT OUTER JOIN

To include all customers, regardless of whether they have orders, you can use a left outer join via `LEFT JOIN` (refer to slides).

If any customers do not have any orders, they will get one record where the `CUSTOMER_ORDER` fields will be null.

```
SELECT CUSTOMER.CUSTOMER_ID,  
NAME,  
STREET_ADDRESS,  
CITY,  
STATE,  
ZIP,  
ORDER_DATE,  
SHIP_DATE,  
ORDER_ID,  
PRODUCT_ID,
```

```
ORDER_QTY
```

```
FROM CUSTOMER LEFT JOIN CUSTOMER_ORDER  
ON CUSTOMER.CUSTOMER_ID = CUSTOMER_ORDER.CUSTOMER_ID
```

7.3 Finding Customers with No Orders

With a left outer join, you can filter for NULL values on the CUSTOMER_ORDER table to find customers that have no orders.

```
SELECT CUSTOMER.CUSTOMER_ID,  
NAME AS CUSTOMER_NAME  
  
FROM CUSTOMER LEFT JOIN CUSTOMER_ORDER  
ON CUSTOMER.CUSTOMER_ID = CUSTOMER_ORDER.CUSTOMER_ID  
  
WHERE ORDER_ID IS NULL
```

You can use a left outer join to find child records with no parent, or parent records with no children (e.g. a CUSTOMER_ORDER with no CUSTOMER, or a CUSTOMER with no CUSTOMER_ORDERS).

7.4 Joining Multiple Tables

Bring in PRODUCT to supply product information for each CUSTOMER_ORDER, on top of CUSTOMER information.

```
SELECT ORDER_ID,  
CUSTOMER.CUSTOMER_ID,  
NAME AS CUSTOMER_NAME,  
STREET_ADDRESS,  
CITY,  
STATE,  
ZIP,  
ORDER_DATE,  
PRODUCT.PRODUCT_ID,  
DESCRIPTION,  
ORDER_QTY  
  
FROM CUSTOMER INNER JOIN CUSTOMER_ORDER  
ON CUSTOMER.CUSTOMER_ID = CUSTOMER_ORDER.CUSTOMER_ID  
  
INNER JOIN PRODUCT  
ON CUSTOMER_ORDER.PRODUCT_ID = PRODUCT.PRODUCT_ID
```

7.7 Using Expressions with JOINS

You can use expressions combining any fields on any of the joined tables. For instance, we can now get the total revenue for each customer.

```
SELECT ORDER_ID,  
CUSTOMER.CUSTOMER_ID,  
NAME AS CUSTOMER_NAME,  
STREET_ADDRESS,  
CITY,  
STATE,  
ZIP,  
ORDER_DATE,  
PRODUCT.PRODUCT_ID,  
DESCRIPTION,  
ORDER_QTY,  
ORDER_QTY * PRICE as REVENUE  
  
FROM CUSTOMER INNER JOIN CUSTOMER_ORDER  
ON CUSTOMER.CUSTOMER_ID = CUSTOMER_ORDER.CUSTOMER_ID  
  
INNER JOIN PRODUCT  
ON CUSTOMER_ORDER.PRODUCT_ID = PRODUCT.PRODUCT_ID
```

7.6 Using GROUP BY with JOINS

You can use GROUP BY with a join. For instance, you can find the total revenue for each customer by leveraging all three joined tables, and aggregating the REVENUE expression we created earlier.

```
SELECT  
CUSTOMER.CUSTOMER_ID,  
NAME AS CUSTOMER_NAME,  
sum(ORDER_QTY * PRICE) as TOTAL_REVENUE  
  
FROM CUSTOMER INNER JOIN CUSTOMER_ORDER  
ON CUSTOMER.CUSTOMER_ID = CUSTOMER_ORDER.CUSTOMER_ID  
  
INNER JOIN PRODUCT  
ON CUSTOMER_ORDER.PRODUCT_ID = PRODUCT.PRODUCT_ID  
  
GROUP BY 1,2
```

To see all customers even if they had no orders, use a LEFT JOIN

```
SELECT  
CUSTOMER.CUSTOMER_ID,
```



```

NAME AS CUSTOMER_NAME,
sum(ORDER_QTY * PRICE) as TOTAL_REVENUE

FROM CUSTOMER LEFT JOIN CUSTOMER_ORDER
ON CUSTOMER.CUSTOMER_ID = CUSTOMER_ORDER.CUSTOMER_ID

LEFT JOIN PRODUCT
ON CUSTOMER_ORDER.PRODUCT_ID = PRODUCT.PRODUCT_ID

GROUP BY 1,2

```

You can also use a `coalesce()` function to turn null sums into zeros.

```

SELECT
CUSTOMER.CUSTOMER_ID,
NAME AS CUSTOMER_NAME,
coalesce(sum(ORDER_QTY * PRICE), 0) as TOTAL_REVENUE

FROM CUSTOMER LEFT JOIN CUSTOMER_ORDER
ON CUSTOMER.CUSTOMER_ID = CUSTOMER_ORDER.CUSTOMER_ID

LEFT JOIN PRODUCT
ON CUSTOMER_ORDER.PRODUCT_ID = PRODUCT.PRODUCT_ID

GROUP BY 1,2

```

Section VIII - Database Design

Refer to slides for database design concepts

To view source code for SQL Injection Demo, here is the GitHub page:
<https://github.com/thomasniel/sql-injection-demo>

To read about normalized forms (which we do not cover in favor of a more intuitive approach), you can read this article:

<http://www.dummies.com/programming/sql/sql-first-second-and-third-normal-forms/>

7.1 - Creating a Table

In SQLiteStudio, navigate to *Database* -> *Add a Database* and click the green plus icon to create a new database. Choose a location and name it `surgetech_conference.db`.

Create the COMPANY table. To create a new table, use the SQLiteStudio wizard by right-clicking the `surgetech_conference` database and selecting **Create a table**. You can also just execute the following SQL.

```
CREATE TABLE COMPANY (  
    COMPANY_ID INTEGER PRIMARY KEY AUTOINCREMENT,  
    NAME VARCHAR(30) NOT NULL,  
    DESCRIPTION VARCHAR(60),  
    PRIMARY_CONTACT_ATTENDEE_ID INTEGER NOT NULL,  
    FOREIGN KEY (PRIMARY_CONTACT_ATTENDEE_ID) REFERENCES ATTENDEE(ATTENDEE_ID)  
);
```

After each field declaration, we create “rules” for that field. For example, `COMPANY_ID` must be an `INTEGER`, it is a `PRIMARY KEY`, and it will `AUTOINCREMENT` to automatically generate a consecutive integer ID for each new record. The `NAME` field holds text because it is `VARCHAR` (a variable number of characters), and it is limited to 30 characters and cannot be `NULL`.

Lastly, we declare any `FOREIGN KEY` constraints, specifying which field is a `FOREIGN KEY` and what `PRIMARY KEY` it references. In this example, `PRIMARY_CONTACT_ATTENDEE_ID` “references” the `ATTENDEE_ID` in the `ATTENDEE` table, and it can only be those values.

7.2 - Creating the other tables

Create the other tables using the SQLiteStudio *New table* wizard, or just executing the following SQL code.

```
CREATE TABLE ROOM (  
    ROOM_ID INTEGER PRIMARY KEY AUTOINCREMENT,  
    FLOOR_NUMBER INTEGER NOT NULL,  
    SEAT_CAPACITY INTEGER NOT NULL  
);  
  
CREATE TABLE PRESENTATION (  
    PRESENTATION_ID INTEGER PRIMARY KEY AUTOINCREMENT,  
    BOOKED_COMPANY_ID INTEGER NOT NULL,  
    BOOKED_ROOM_ID INTEGER NOT NULL,  
    START_TIME TIME,  
    END_TIME TIME,  
    FOREIGN KEY (BOOKED_COMPANY_ID) REFERENCES COMPANY(COMPANY_ID)  
    FOREIGN KEY (BOOKED_ROOM_ID) REFERENCES ROOM(ROOM_ID)  
);  
  
CREATE TABLE ATTENDEE (  
    ATTENDEE_ID INTEGER PRIMARY KEY AUTOINCREMENT,  
    FIRST_NAME VARCHAR (30) NOT NULL,
```

```

    LAST_NAME VARCHAR (30) NOT NULL,
    PHONE INTEGER,
    EMAIL VARCHAR (30),
    VIP BOOLEAN DEFAULT (0)
);

CREATE TABLE PRESENTATION_ATTENDANCE (
    TICKET_ID INTEGER PRIMARY KEY AUTOINCREMENT,
    PRESENTATION_ID INTEGER,
    ATTENDEE_ID INTEGER,
    FOREIGN KEY (PRESENTATION_ID) REFERENCES PRESENTATION(PRESENTATION_ID)
    FOREIGN KEY (ATTENDEE_ID) REFERENCES ATTENDEE(ATTENDEE_ID)
);

```

Creating Views

It is not uncommon to save `SELECT` queries that are used frequently into a database. These are known as **Views** and act very similarly to tables. You can essentially save a `SELECT` query and work with it just like a table.

For instance, say we wanted to save this SQL query that includes `ROOM` and `COMPANY` info with each `PRESENTATION` record.

```

SELECT COMPANY.NAME as BOOKED_COMPANY,
ROOM.ROOM_ID as ROOM_NUMBER,
ROOM.FLOOR_NUMBER as FLOOR,
ROOM.SEAT_CAPACITY as SEATS,
START_TIME, END_TIME

FROM PRESENTATION

INNER JOIN COMPANY
ON PRESENTATION.BOOKED_COMPANY_ID = COMPANY.COMPANY_ID

INNER JOIN ROOM
ON PRESENTATION.BOOKED_ROOM_ID = ROOM.ROOM_ID

```

You can save this as a view by right-clicking *Views* in the database navigator, and then *Create a view*. You can then paste the SQL as the body and give the view a name, such as `PRESENTATION_VW` (where “VW” means “View”).

You can also just execute the following SQL syntax: `CREATE [view name] AS [a SELECT query]`. For this example, this is what it would look like.

```

CREATE VIEW PRESENTATION_VW AS

SELECT COMPANY.NAME as BOOKED_COMPANY,

```

```

ROOM.ROOM_ID as ROOM_NUMBER,
ROOM.FLOOR_NUMBER as FLOOR,
ROOM.SEAT_CAPACITY as SEATS,
START_TIME, END_TIME

FROM PRESENTATION

INNER JOIN COMPANY
ON PRESENTATION.BOOKED_COMPANY_ID = COMPANY.COMPANY_ID

INNER JOIN ROOM
ON PRESENTATION.BOOKED_ROOM_ID = ROOM.ROOM_ID

```

You will then see the PRESENTATION_VW in your database navigator, and you can query it just like a table.

```

SELECT * FROM PRESENTATION_VW
WHERE SEATS >= 30

```

Obviously, there is no data yet so you will not get any results. But there will be once you populate data into this database.

Section IX - Writing Data

In this section, we will learn how to write, modify, and delete data in a database.

9.1 Using INSERT

To create a new record in a table, use the INSERT command and supply the values for the needed columns.

Put yourself into the ATTENDEE table.

```

INSERT INTO ATTENDEE (FIRST_NAME, LAST_NAME)
VALUES ('Thomas', 'Niield')

```

Notice above that we declare the table we are writing to, which is ATTENDEE. Then we declare the columns we are supplying values for (FIRST_NAME, LAST_NAME), followed by the values for this new record ('Thomas', 'Niield').

Notice we did not have to supply a value for ATTENDEE_ID as we have set it in the previous section to generate its own value. PHONE, EMAIL, and VIP fields have default values or are nullable, and therefore optional.

9.2 Multiple INSERT records

You can insert multiple rows in an INSERT. This will add three people to the ATTENDEE table.

```
INSERT INTO ATTENDEE (FIRST_NAME, LAST_NAME, PHONE, EMAIL, VIP)
VALUES ('Jon', 'Skeeter', 4802185842, 'john.skeeter@rex.net', 1),
       ('Sam', 'Scala', 2156783401, 'sam.scala@gmail.com', 0),
       ('Brittany', 'Fisher', 5932857296, 'brittany.fisher@outlook.com', 0)
```

9.3 Testing the foreign keys

Let's test our design and make sure our primary/foreign keys are working.

Try to INSERT a COMPANY with a PRIMARY_CONTACT_ATTENDEE_ID that does not exist in the ATTENDEE table.

```
INSERT INTO COMPANY (NAME, DESCRIPTION, PRIMARY_CONTACT_ATTENDEE_ID)
VALUES ('RexApp Solutions', 'A mobile app delivery service', 5)
```

Currently, there is no ATTENDEE with an ATTENDEE_ID of 5, this should error out which is good. It means we kept bad data out.

If you use an ATTENDEE_ID value that does exist and supply it as a PRIMARY_CONTACT_ATTENDEE_ID, we should be good to go.

```
INSERT INTO COMPANY (NAME, DESCRIPTION, PRIMARY_CONTACT_ATTENDEE_ID)
VALUES ('RexApp Solutions', 'A mobile app delivery service', 3)
```

9.3 DELETE records

The DELETE command is dangerously simple. To delete records from both the COMPANY and ATTENDEE tables, execute the following SQL commands.

```
DELETE FROM COMPANY;
DELETE FROM ATTENDEE;
```

Note that the COMPANY table has a foreign key relationship with the ATTENDEE table. Therefore we will have to delete records from COMPANY first before it allows us to delete data from ATTENDEE. Otherwise we will get a “FOREIGN KEY constraint failed effort” due to the COMPANY record we just added which is tied to the ATTENDEE with the ATTENDEE_ID of 3.

You can also use a WHERE to only delete records that meet a conditional. To delete all ATTENDEE records with no PHONE or EMAIL, you can run this command.

```
DELETE FROM ATTENDEE
WHERE PHONE IS NULL AND EMAIL IS NULL
```

A good practice is to use a `SELECT *` in place of the `DELETE` first. That way you can get a preview of what records will be deleted with that `WHERE` condition.

```
SELECT * FROM ATTENDEE
WHERE PHONE IS NULL AND EMAIL IS NULL
```

UPDATE records

Say we wanted to change the phone number for the `ATTENDEE` with the `ATTENDEE_ID` value of 3, which is Sam Scala. We can do this with an `UPDATE` statement.

```
UPDATE ATTENDEE SET PHONE = 4802735872
WHERE ATTENDEE_ID = 3
```

Using a `WHERE` is important, otherwise it will update all records with the specified `SET` assignment. This can be handy if you wanted to say, make all `EMAIL` values uppercase.

```
UPDATE ATTENDEE SET EMAIL = UPPER(EMAIL)
```

9.4 Dropping Tables

If you want to delete a table, it also is dangerously simple. Be very careful and sure before you delete any table, because it will remove it permanently.

```
DROP TABLE MY_UNWANTED_TABLE
```

9.5 Transactions

Transactions are helpful when you want a series of writes to succeed.

Below, we execute two successful write operations within a transaction.

```
BEGIN TRANSACTION;

INSERT INTO ROOM (FLOOR_NUMBER, SEAT_CAPACITY) VALUES (9, 80);
INSERT INTO ROOM (FLOOR_NUMBER, SEAT_CAPACITY) VALUES (10, 110);

END TRANSACTION;
```

But if we ever encountered a failure with our write operations, we can call `ROLLBACK` instead of `END TRANSACTION` to go back to the database state when `BEGIN TRANSACTION` was called.

Below, we have a failed operation due to a broken `INSERT`.

```
BEGIN TRANSACTION;
```

```
INSERT INTO ROOM (FLOOR_NUMBER, SEAT_CAPACITY) VALUES (12, 210);  
INSERT INTO ROOM (FLOOR_NUMBER, SEAT_CAPACITY) VALUES (13); --failure
```

So we can call ROLLBACK to “rewind” to the database state when BEGIN TRANSACTION was called.

```
ROLLBACK;
```

9.6 Creating Indexes

You can create an index on a certain column to speed up SELECT performance, such as the price column on the PRODUCT table.

```
CREATE INDEX price_index ON PRODUCT(price);
```

You can also create an index for a column that has unique values, and it will make a special optimization for that case.

```
CREATE UNIQUE INDEX name_index ON CUSTOMER(name);
```

To remove an index, use the DROP command.

```
DROP INDEX price_index;
```

9.7 Working with Dates and Times

Use the ISO ‘yyyy-mm-dd’ syntax with strings to treat them as dates easily.

Keep in mind much of this functionality is proprietary to SQLite. Make sure you learn the date and time functionality for your specific database platform.

```
SELECT * FROM CUSTOMER_ORDER  
WHERE SHIP_DATE < '2015-05-21'
```

To get today’s date:

```
SELECT DATE('now')
```

To shift a date:

```
SELECT DATE('now', '-1 day')  
SELECT DATE('2015-12-07', '+3 month', '-1 day')
```

To work with times, use hh:mm:ss format.

```
SELECT '16:31' < '08:31'
```

To get today’s GMT time:

```
SELECT TIME('now')
```

To shift a time:

```
SELECT TIME('16:31','+1 minute')
```

To merge a date and time, use a DateTime type.

```
SELECT '2015-12-13 16:04:11'
```

```
SELECT DATETIME('2015-12-13 16:04:11','-1 day','+3 hour')
```

To format dates and times a certain way:

```
“sql SELECT strftime('%d-%m-%Y', 'now') ““
```

Refer to SQLite documentation http://www.sqlite.org/lang_datefunc.html

Another helpful tutorial on using dates and times with SQLite. https://www.tutorialspoint.com/sqlite/sqlite_date.htm

Section X - Moving Forward

SQL Resources

Getting Started with SQL (O'Reilly) by Thomas Nield

Learning SQL (O'Reilly) by Alan Beaulieu

Using SQLite (O'Reilly) by Jay A. Kreibich

SQL Practice Problems by Sylvia Moestl Vasilik